


# CNN for Cifar-10 e Cifar-100 Classification on *Coral Edge TPU USB Accelerator* and *ST Cloud JAM*

March 2020


# CNN for Cifar-10 classification on computer

## Neural Network architecture




# CNN for Cifar-10 classification on computer

## Neural Network architecture



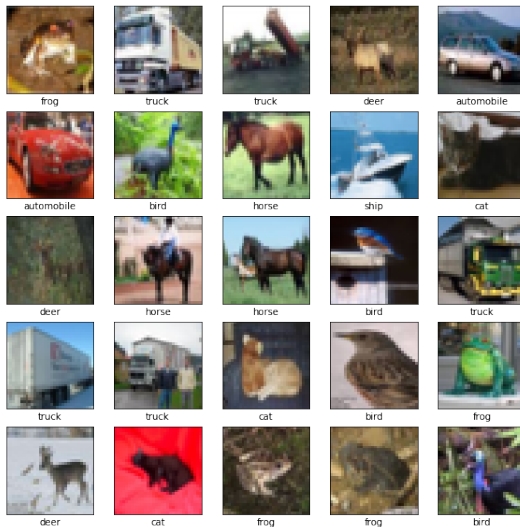
# CNN for Cifar-10 classification on computer

## Neural Network architecture



# Cifar-10 random images

Stampa di pictures random



# CNN for Cifar-10 classification on computer

## Data Augmentation

```
datagen = ImageDataGenerator(  
    ^^Irotation_range=15,  
    ^^Iwidth_shift_range=.1,  
    ^^Iheight_shift_range=.1,  
    ^^Ihorizontal_flip=True,  
    ^^Izoom_range=0.5)
```


# CNN for Cifar-10 classification on computer

## Training hyperparameters

```
model.compile(loss=keras.losses.categorical_crossentropy,  
              optimizer=keras.optimizers.Adam(), metrics=['accuracy'])  
^^I
```

# CNN for Cifar-10 classification on computer

## Accuracy




Epoch 100/100

```
782/782 [=====] - 31s 40ms/step -  
    loss: 0.5471 - acc: 0.8146 - val_loss: 0.3975 - val_acc:  
    0.8701
```




# CNN for Cifar-10 classification on ST Cloud JAM

## Neural Network architecture




# CNN for Cifar-10 classification on ST Cloud JAM

## Neural Network architecture




# CNN for Cifar-10 classification on ST Cloud JAM

## Neural Network architecture



# CNN for Cifar-10 classification on ST Cloud JAM

## Accuracy



```
^^IEpoch 100/100
^^I782/782 [=====] - 43s 55ms/step
    - loss: 0.8155 - acc: 0.7220 - val_loss: 0.5932 -
    val_acc: 0.7963
^^I
```

# Test CNN con Cifar-10 su ST Cloud JAM

## Validation on target of the Neural Network

The network is imported without weights compression. The L2R error between the results on host (computer) and target is in the order of  $10^{-6}$

```
acc=100.00%, rmse=0.0000, mae=0.0000
```

```
10 classes (1000 samples)
```


```
-----
C0      0      .      .      .      .      .      .      .      .
C1      .    508      .      .      .      .      .      .      .
C2      .      .      0      .      .      .      .      .      .
C3      .      .      .      0      .      .      .      .      .
C4      .      .      .      .      0      .      .      .      .
C5      .      .      .      .      .    14      .      .      .
C6      .      .      .      .      .      .      0      .      .
C7      .      .      .      .      .      .      .      5      .
C8      .      .      .      .      .      .      .      .      2
C9      .      .      .      .      .      .      .      .    471
```

```
Evaluation report (summary)
```

```
-----
Mode                acc                rmse                mae
-----
X-cross #1          100.0%          0.000001    0.000000
L2r error : 4.15037266e-06 (expected to be < 0.01)
```


# CNN for Cifar-100 classification on computer

## Neural Network architecture




# CNN for Cifar-100 classification on computer

## Neural Network architecture




# CNN for Cifar-100 classification on computer

## Neural Network architecture






# Cifar-100 random images



# CNN for Cifar-100 classification on computer

## Accuracy



Epoch 120/120

```
782/782 [=====] - 50s 64ms/step -  
    loss: 0.6304 - acc: 0.8049 - val_loss: 1.2676 - val_acc:  
    0.7001
```

# Test CNN con Cifar-100 su ST Cloud JAM

## Validation on target of the Neural Network

The network is imported without weights compression. The L2R error between the results on host (computer) and target is in the order of  $10^{-6}$

Running with inputs (1000, 32, 32, 3) ..

```
acc=100.00%, rmse=0.0000, mae=0.0000
```

```
Evaluation report (summary)
```

Mode	acc	rmse	mae
X-cross #1	100.0%	0.000000	0.000000

```
L2r error : 1.32934053e-06 (expected to be < 0.01)
```

# CNN for Cifar-100 classification on Coral Edge TPU

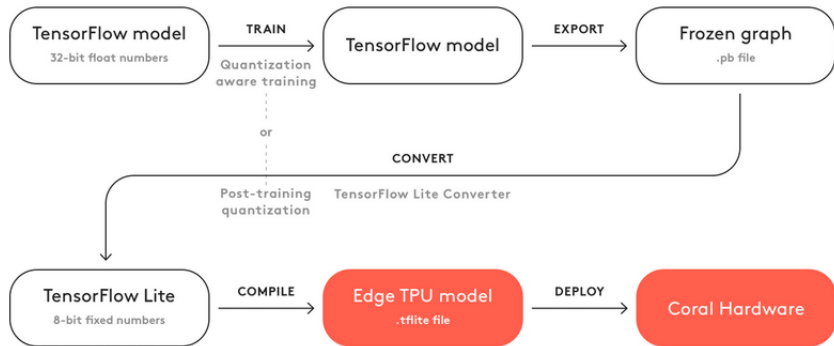
## Introduction

Coral USB Accelerator is a device aimed to support a calculator with an operating system.



# CNN for Cifar-100 classification on Coral Edge TPU

## Requirement, Structure



# CNN for Cifar-100 classification on Coral Edge TPU

## Requirement, Structure

```
def cnn_model_fn(features, labels, mode):  
    with tf.name_scope('model_input') as scope:  
        input_layer = tf.reshape(features, [-1, 32, 32, 3], name  
="input")  
  
        with tf.name_scope('model_conv1') as scope:  
            conv1 = tf.layers.conv2d(inputs=input_layer, filters=32,  
                kernel_size=[2, 2],  
                padding="same", activation=tf.nn.relu6,  
                trainable=mode == tf.estimator.ModeKeys.TRAIN)  
  
            with tf.name_scope('model_conv2') as scope:  
                conv2 = tf.layers.conv2d(inputs=conv1, filters=32,  
                    kernel_size=[2, 2],  
                    padding="same", activation=tf.nn.relu6,  
                    trainable=mode == tf.estimator.ModeKeys.TRAIN)  
  
                pool1 = tf.layers.max_pooling2d(inputs=conv2, pool_size  
=[2, 2], strides=(2, 2))  
                drop1 = tf.layers.dropout(inputs=pool1, rate=0.2)
```

# CNN for Cifar-100 classification on Coral Edge TPU

## Requirement, Structure

```
with tf.name_scope('model_conv3') as scope:
    conv3 = tf.layers.conv2d(inputs=drop1, filters=64,
kernel_size=[2, 2],
padding="same", activation=tf.nn.relu6,
trainable=mode == tf.estimator.ModeKeys.TRAIN)

with tf.name_scope('model_conv4') as scope:
    conv4 = tf.layers.conv2d(inputs=conv3, filters=64,
kernel_size=[2, 2],
padding="same", activation=tf.nn.relu6,
trainable=mode == tf.estimator.ModeKeys.TRAIN)
    pool2 = tf.layers.max_pooling2d(inputs=conv4, pool_size
=[2, 2], strides=(2, 2))
    drop2 = tf.layers.dropout(inputs=pool2, rate=0.3)
```

# CNN for Cifar-100 classification on Coral Edge TPU

## Tensorflow network definition

```
with tf.name_scope('model_conv5') as scope:
    conv5 = tf.layers.conv2d(inputs=drop2, filters=128,
kernel_size=[2, 2],
padding="same", activation=tf.nn.relu6,
trainable=mode == tf.estimator.ModeKeys.TRAIN)

with tf.name_scope('model_conv6') as scope:
    conv6 = tf.layers.conv2d(inputs=conv5, filters=128,
kernel_size=[2, 2],
padding="same", activation=tf.nn.relu6,
trainable=mode == tf.estimator.ModeKeys.TRAIN)
    pool3 = tf.layers.max_pooling2d(inputs=conv6, pool_size
=[2, 2], strides=(2, 2))
    drop3 = tf.layers.dropout(inputs=pool3, rate=0.4)
```



# CNN for Cifar-100 classification on Coral Edge TPU

## Tensorflow network definition

```
with tf.name_scope('model_dense') as scope:
    flat = tf.reshape(drop3, [-1, 4 * 4 * 128])

    dense = tf.layers.dense(inputs=flat, units=1024,
        activation=tf.nn.relu6,
        trainable=mode == tf.estimator.ModeKeys.TRAIN)

    drop4 = tf.layers.dropout(inputs=dense, rate=0.5,
        training=mode == tf.estimator.ModeKeys.TRAIN)

    with tf.name_scope('model_output') as scope:
        logits = tf.layers.dense(inputs=drop4, units=10,
            trainable=mode == tf.estimator.ModeKeys.TRAIN)

        predictions = {latin1
            "classes": tf.argmax(input=latin1=logits, axis=1),
            "probabilities": tf.nn.softmax(logits, name="
softmax_tensor")}
}
```

# CNN for Cifar-100 classification on Coral Edge TPU

## Tensorflow network definition

```
# TRAIN mode
if mode == tf.estimator.ModeKeys.TRAIN:
    g = tf.get_default_graph()
    tf.contrib.quantize.create_training_graph(input_graph=g,
        quant_delay=6000)
    optimizer = tf.train.AdamOptimizer(learning_rate=0.001)
    train_op = optimizer.minimize(
        loss=loss,
        global_step=tf.train.get_global_step())
    return tf.estimator.EstimatorSpec(mode=mode, loss=loss,
        train_op=train_op)

# EVAL mode
g = tf.get_default_graph()
tf.contrib.quantize.create_eval_graph(input_graph=g)
labels = tf.argmax(labels, axis=1)
eval_metric_ops = {
    "accuracy": tf.metrics.accuracy(labels=labels,
        predictions=predictions["classes"])
}
return tf.estimator.EstimatorSpec(
    mode=mode, loss=loss, eval_metric_ops=eval_metric_ops)
```

# CNN for Cifar-100 classification on Coral Edge TPU

## Training and saving

### Training:

```
def fit_all_batches(xEpochs):  
    # Train the model  
    train_input_fn = tf.estimator.inputs.numpy_input_fn(  
        x=train_features,  
        y=train_labels,  
        batch_size=32,  
        num_epochs=xEpochs,  
        shuffle=False)  
  
    # train one step and display the probabilities  
    cifar10_classifier.train(  
        input_fn=train_input_fn,  
        steps=None,  
        hooks=[logging_hook])  
    ^^I  
fit_all_batches(100)
```

### Saving in .pb format:


```
cifar10_classifier.experimental_export_all_saved_models(  
    savedir,  
    input_receiver_fn_map)
```

# CNN for Cifar-100 classification on Coral Edge TPU

## Accuracy and test

```
{'accuracy': 0.7656, 'loss': 1.5887853, 'global_step': 142032}
```

```
[['airplane' '0.9999999']  
 ['automobile' '1.5820127e-23']  
 ['bird' '6.063283e-08']  
 ['cat' '7.3694817e-10']  
 ['deer' '2.7018123e-08']  
 ['dog' '1.1388865e-14']  
 ['frog' '5.9821274e-15']  
 ['horse' '1.180477e-18']  
 ['ship' '8.983325e-14']  
 ['truck' '2.4614857e-21']]
```



# CNN for Cifar-100 classification on Coral Edge TPU

## Frozen graph creation

script Python freeze\_graph.py

```
tf.saved_model.loader.load(sess, ["serve"], savedir )

from tensorflow.python.tools import freeze_graph
from tensorflow.python.saved_model import tag_constants

input_graph_filename = None
input_saved_model_dir = savedir
output_node_names = "softmax_tensor"
output_graph_filename = os.path.join(savedir, "frozen.pb")
input_binary = False
input_saver_def_path = False
restore_op_name = None
filename_tensor_name = None
clear_devices = True
input_meta_graph = False
checkpoint_path = None
```

# CNN for Cifar-100 classification on Coral Edge TPU

## TFlite conversion

**TOCO** converter, disponibile con Tensorflow

```
toco \  
--input_file=frozen.pb \  
--output_file=tflite_model.tflite \  
--input_format=TENSORFLOW_GRAPHDEF \  
--output_format=TFLITE \  
--inference_type=QUANTIZED_UINT8 \  
--input_shape="1,32,32,3" \  
--input_array=model_input/numberstyleinput \  
--output_array=softmax_tensor \  
--std_dev_values=127 \  
--mean_value=127 \  
--default_ranges_min=0 \  
--default_ranges_max=6
```

# CNN for Cifar-100 classification on Coral Edge TPU

## Compilation

### Edge TPU Compiler

→ conversione da TF Lite a in un modello per Coral USB Accelerator

```
edgetpu_compiler [options] model...
```

# CNN for Cifar-100 classification on Coral Edge TPU

## Classification

**Classificatore:** CIFAR10\_Inference\_EdgeTPU.py

```
python3 my_classifier.py --model=/home/matteo/Documents/
    Frozen/Compiled_models/frozen_edgetpu.tflite --image=/
    home/matteo/Downloads/cifar10_images/truck3.png
WARNING:root:From my_classifier.py:30: The name
    ClassifyWithImage will be deprecated. Please use
    classify_with_image instead.

-----
truck
Score :    0.99609375
Elapsed time for the last inference:    0.03292989730834961
```



# CNN for Cifar-100 classification on Coral Edge TPU

## Classification

```
python3 my_classifier.py --model=/home/matteo/Documents/
Frozen/Compiled_models/frozen_example_min_max_edgetpu.
tflite --image=/home/matteo/Downloads/cifar10_images/
bird3.png
WARNING:root:From my_classifier.py:30: The name
ClassifyWithImage will be deprecated. Please use
classify_with_image instead.

-----
truck
Score : 0.62890625
-----
frog
Score : 0.2890625
Elapsed time for the last inference: 0.03208661079406738
```